

Custom Semantic Image Segmentation with DeepLabv3+

Computational Intelligence Lab in the Department of
Computer Science at Cal Poly Pomona

David Hughes
Faculty Mentor: Dr. Hao Ji

Resources

- Deeplabv3+ paper
 - <https://arxiv.org/abs/1802.02611>
- Tensorflow implementation of deeplab
 - <https://github.com/tensorflow/models/tree/master/research/deeplab>
- Github for custom training with Tensorflow's deeplab
 - https://github.com/David-Hughes3/deeplab_project

Semantic Image Segmentation



predict



Person
Bicycle
Background

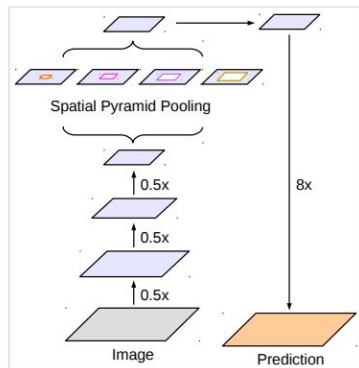
<http://host.robots.ox.ac.uk/pascal/VOC/voc2012/#devkit>

Deeplabv3+

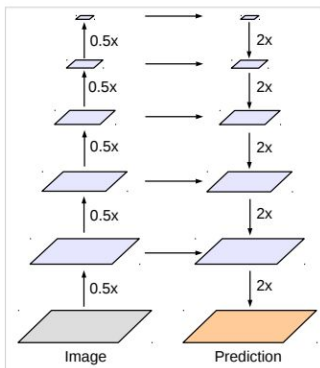
Goal of Deeplabv3+

Combine the advantages of:

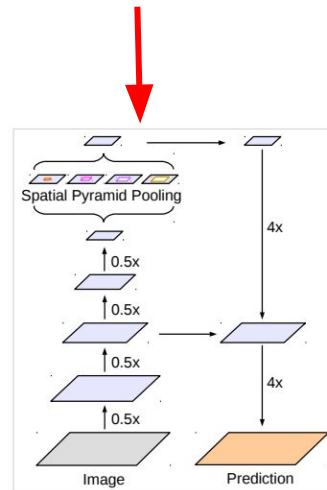
- Spatial pyramid pooling module
 - Encode multi-scale contextual information
 - By probing incoming features with filters or pooling ops at multiple rates and effective FOVs
- Encoder-decoder structure
 - Capture sharper object boundaries
 - By recovering spatial information



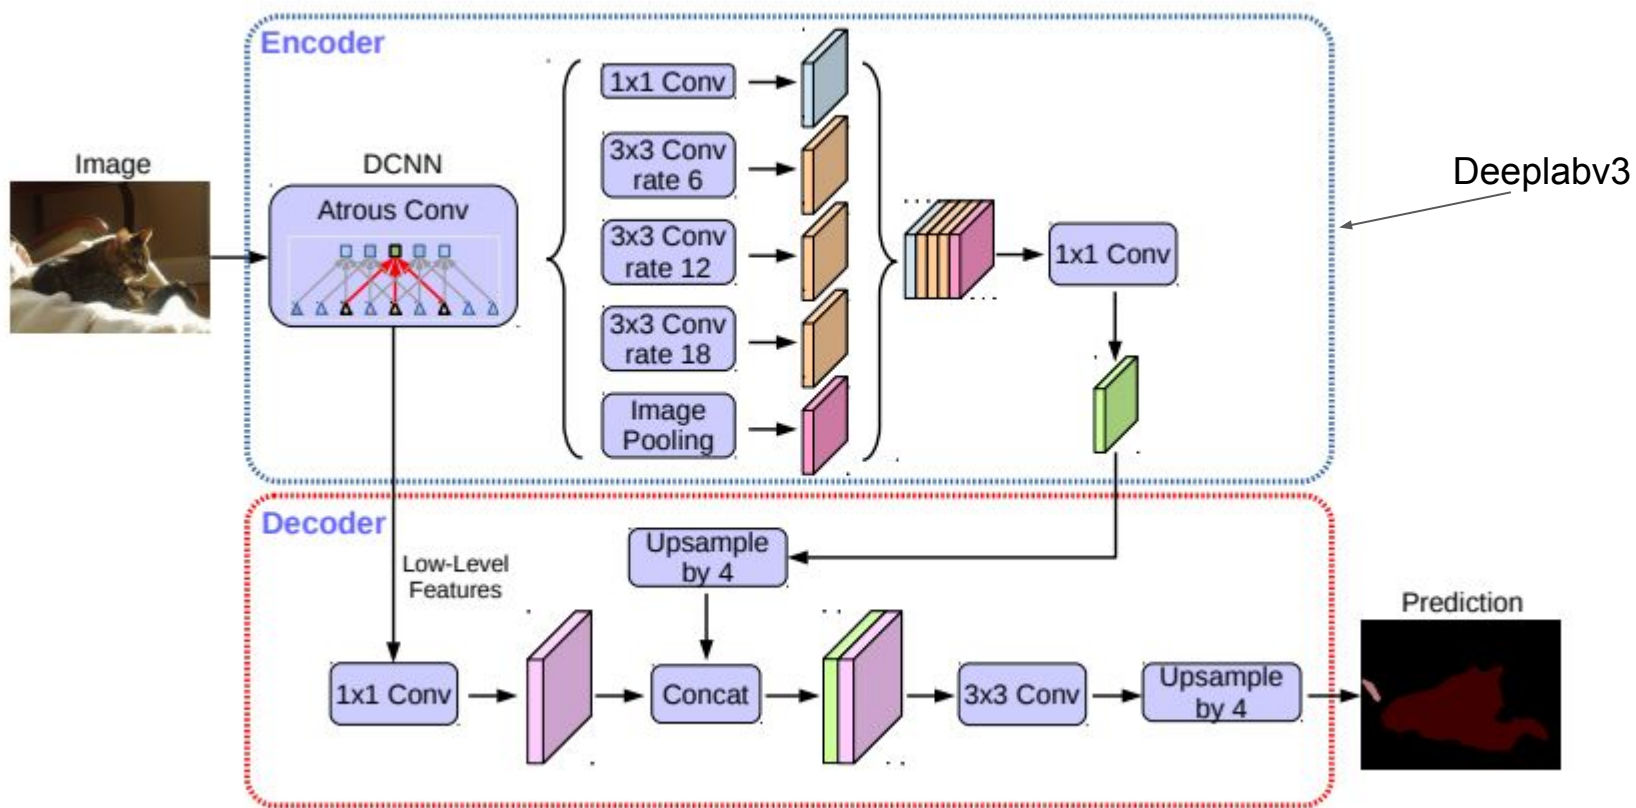
(a) Spatial Pyramid Pooling



(b) Encoder-Decoder



(c) Encoder-Decoder with Atrous Conv

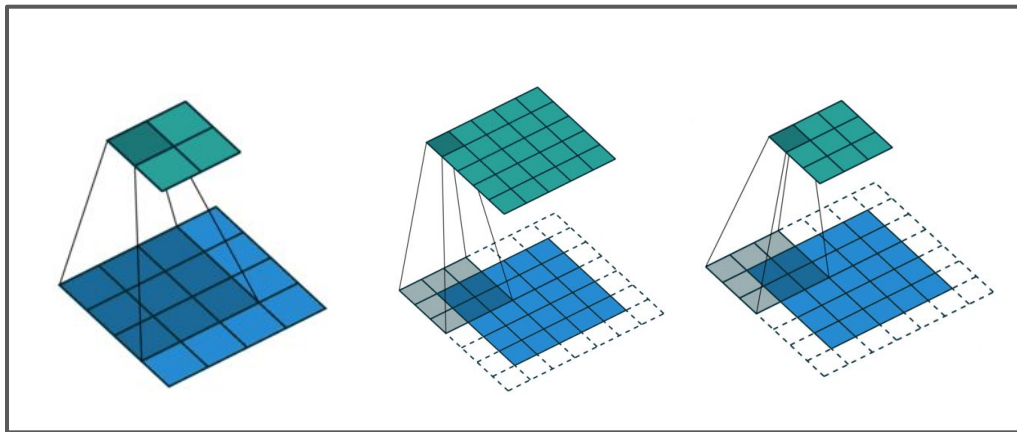


Deeplabv3+ Structure

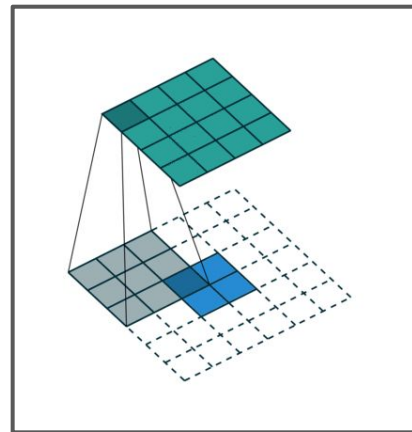
Encoder = DeepLabv3

Basic Types of Convolutions

Blue maps are inputs, and cyan maps are outputs.



Convolution



Transposed
Convolution

Atrous Convolution / Hole Algorithm/ Dilated Conv

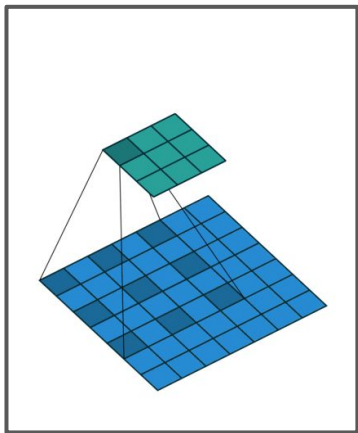
$$y[i] = \sum_k x[i + r \cdot k] w[k]$$

Atrous Convolution

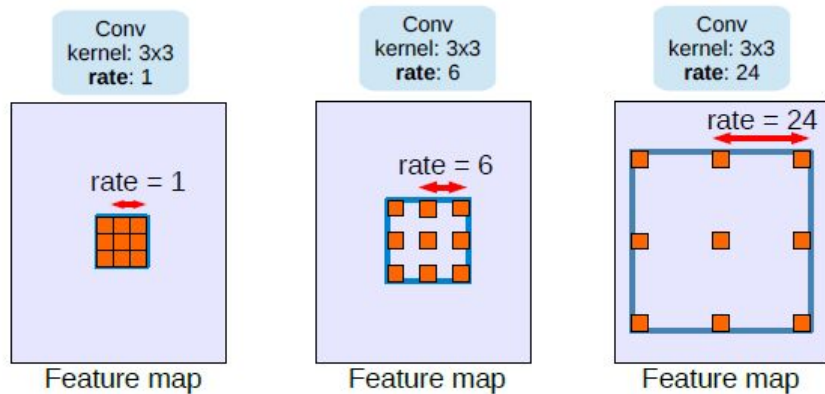
- Location i on output feature map y
- Conv filter w
- Applied over the input feature map x
- Atrous rate r is stride to sample input signal

- $r=1$, standard convolution
- Holes = introducing zeros between filter values
 - “convolving the input x with upsampled filters produced by inserting $r-1$ zeros between two consecutive filter values along each spatial dimension.” [source](#)
- **Filter’s field-of-view can be changed by modifying rate**
 - Small = accurate localization
 - Large = context assimilation

Atrous Convolution



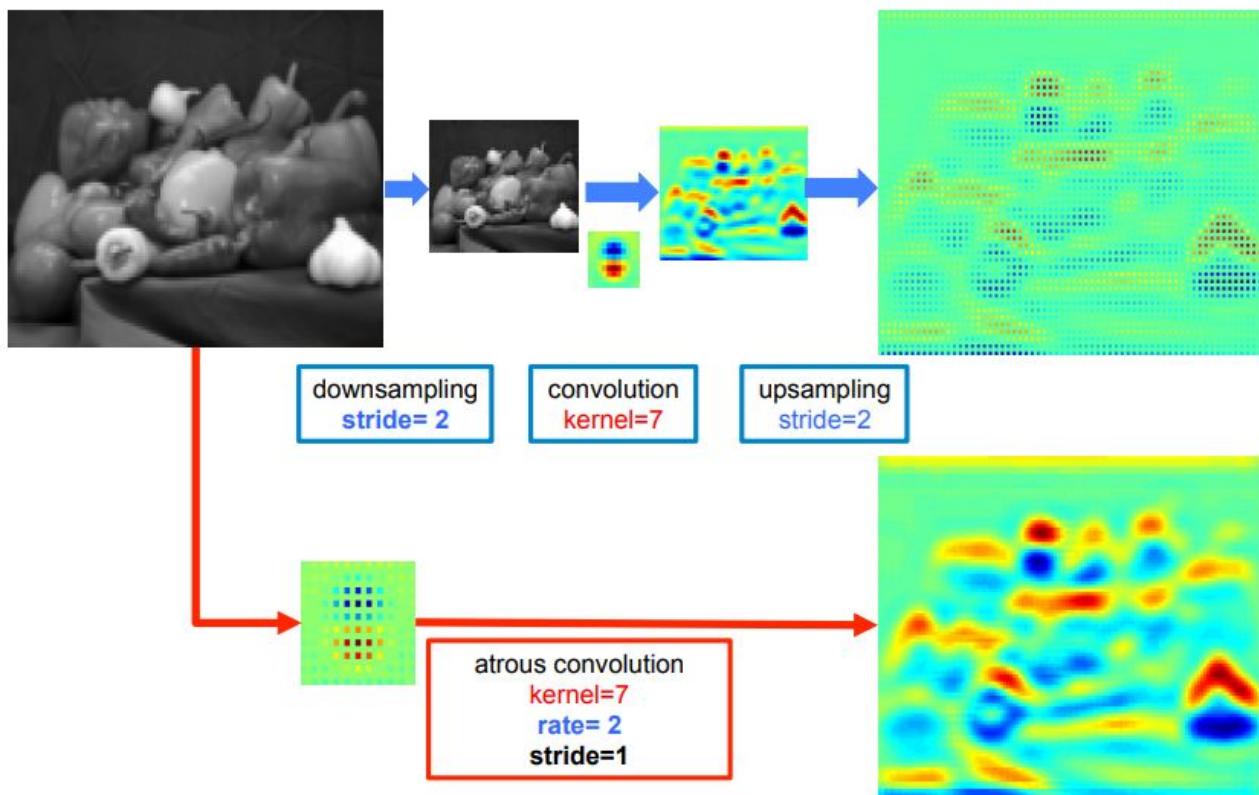
Dilated-
Convolution



Atrous Convolution with Different Rates r

Reasoning

- repeated combination of max-pooling and striding at consecutive layers of these networks reduces significantly the spatial resolution of the resulting feature maps
- A partial remedy is to use 'deconvolutional'/transpose layers
 - However requires additional memory and time.
- Atrous convolution allows computing the responses of any layer at any desirable resolution

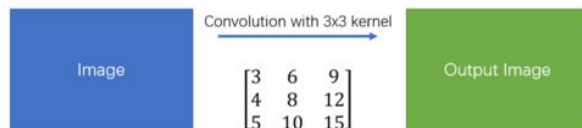


Depthwise separable convolution

First, spatial separable convolutions

- Instead of doing one conv with 9 multiplications
- Do two convolutions with 3 mults each (6 total)
- Less mults = less comp complexity

Simple Convolution



Spatial Separable Convolution




<https://towardsdatascience.com/a-basic-introduction-to-separable-convolutions-b99ec3102728>

Depthwise separable convolution

- Not all kernels can be factored into two smaller kernels
 - Depthwise separable convolution work either way = more common
- Ex: `keras.layers.SeparableConv2D` or `tf.layers.separable_conv2d`
- Like spatial separable conv, depthwise splits a kernel into 2 separate kernels
 - ***Depthwise convolution kernel***
 - ***Pointwise convolution kernel***

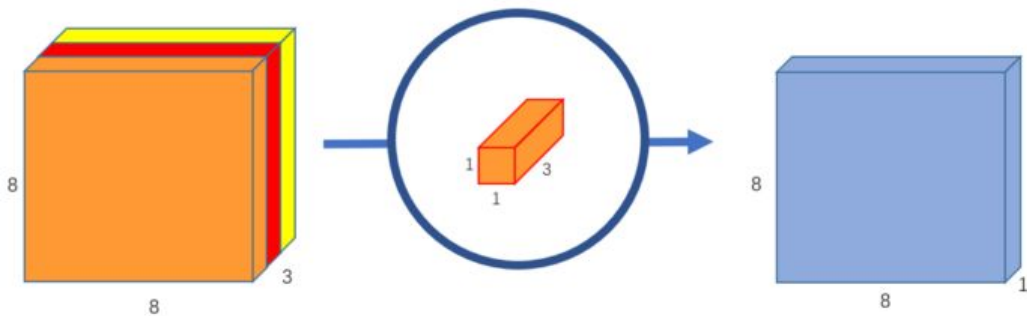
Depthwise Convolution Kernel

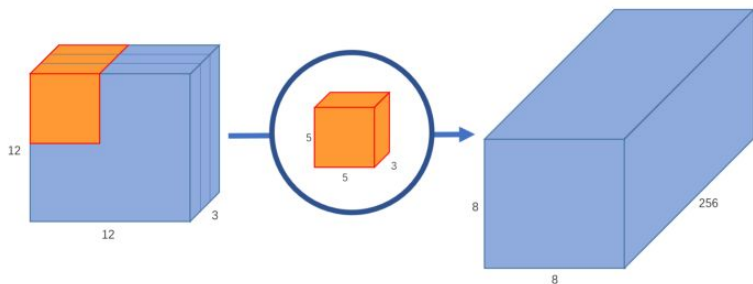
- Performs a spatial convolution independently for each channel
 - Apply a single filter for each input channel



Pointwise convolution kernel


- Combines the output from the depthwise convolution





Normal Convolution:

- 12x12x3 image with 256 (5x5x3) kernels -> 8x8x1 image
- 256 (5x5x3) kernels move 8x8 times
 - $256 \times 3 \times 5 \times 5 \times 8 \times 8 = 1,228,800$ multiplications.



Depthwise Convolution (conv without changing depth):

- 12x12x3 image with 3 (5x5x1) kernels
- 3 (5x5x1) kernels that move 8x8 times.
 - That's $3 \times 5 \times 5 \times 8 \times 8 = 4,800$ multiplications

Pointwise Convolution:

- 256 (1x1x3) kernels that output a 8x8x1 image each to get a final image of shape 8x8x256.
- 256 (1x1x3) kernels that move 8x8 times.
 - That's $256 \times 1 \times 1 \times 3 \times 8 \times 8 = 49,152$ multiplications.

$$49,152 + 4,800 = 53,952 << 1,228,800$$

Atrous Separable Convolution

DeepLabv3+: Encoder-Decoder with Atrous Separable Convolution

5




Fig. 3. 3×3 Depthwise separable convolution decomposes a standard convolution into (a) a depthwise convolution (applying a single filter for each input channel) and (b) a pointwise convolution (combining the outputs from depthwise convolution across channels). In this work, we explore *atrous separable convolution* where atrous convolution is adopted in the depthwise convolution, as shown in (c) with *rate* = 2.

Atrous Spatial Pyramid Pooling (ASPP) module

- probes convolutional features at multiple scales by applying atrous convolution with different rates

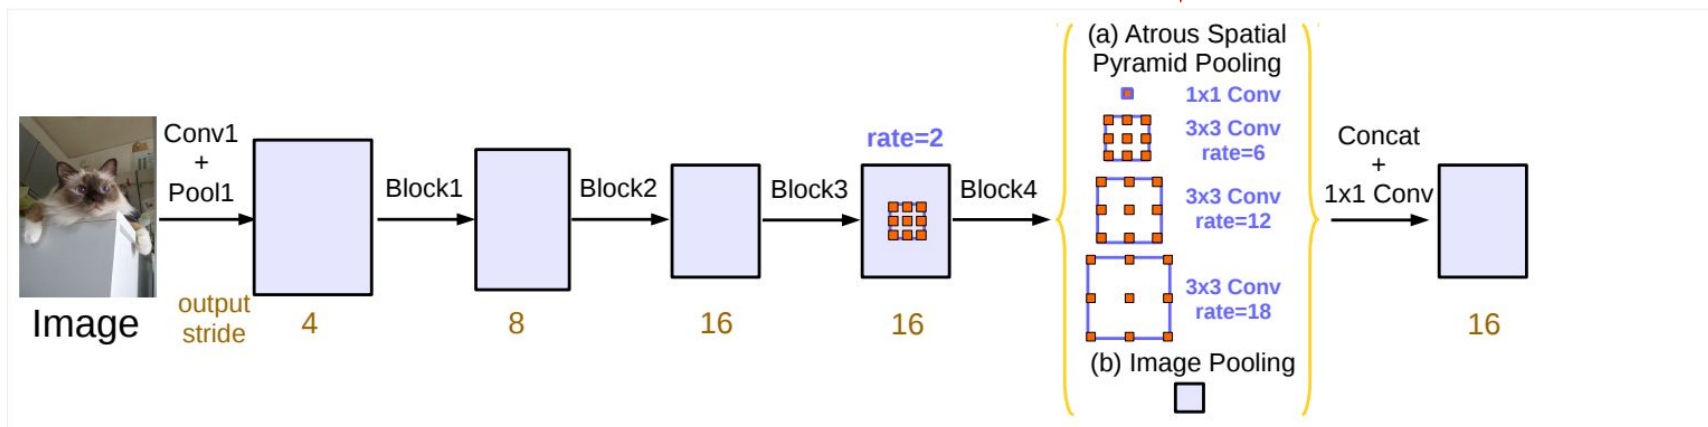
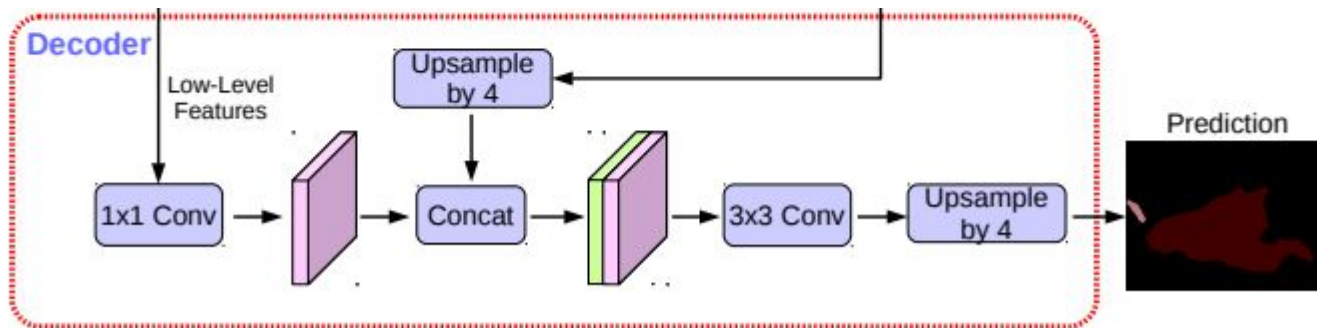


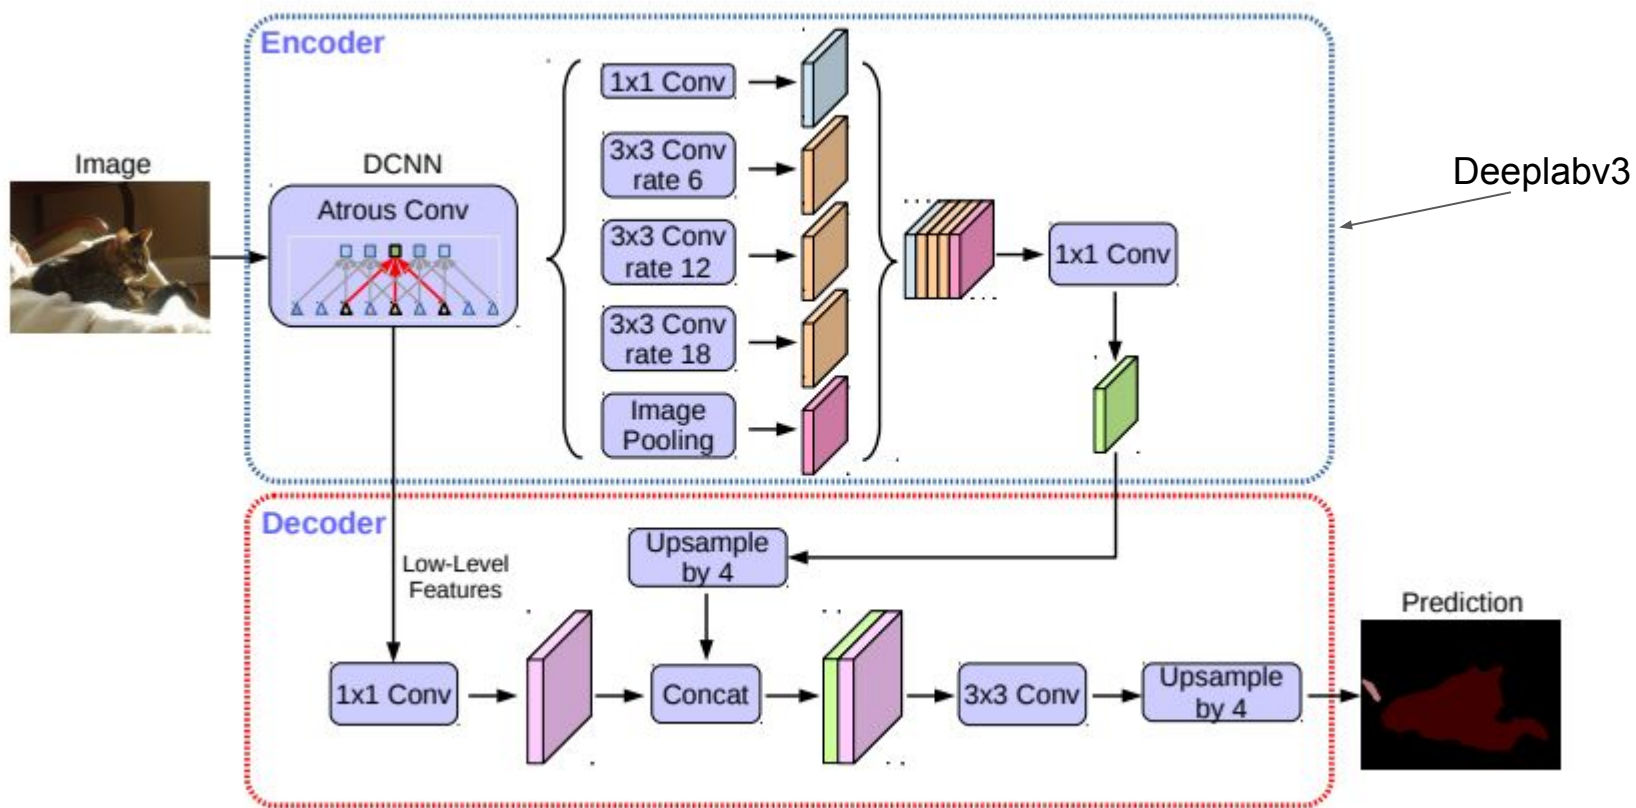
Figure 5. Parallel modules with atrous convolution (ASPP), augmented with image-level features.

Decoder

Simple Decoder

- very simple yet effective decoder module:
 - concatenation of
 - DeepLabv3 feature map
 - channel-reduced Conv2 feature map
 - refined by two $[3 \times 3, 256]$ operations.





Deeplabv3+ Structure

Deeplabv3+ Paper results

Method	mIOU
Deep Layer Cascade (LC) [82]	82.7
TuSimple [77]	83.1
Large_Kernel_Matters [60]	83.6
Multipath-RefineNet [58]	84.2
ResNet-38_MS_COCO [83]	84.9
PSPNet [24]	85.4
IDW-CNN [84]	86.3
CASIA_IVA_SDN [63]	86.6
DIS [85]	86.8
DeepLabv3 [23]	85.7
DeepLabv3-JFT [23]	86.9
DeepLabv3+ (Xception)	87.8
DeepLabv3+ (Xception-JFT)	89.0

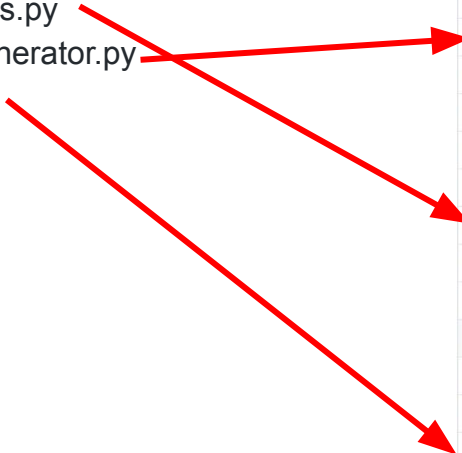
Table 6. PASCAL VOC 2012 *test* set results with top-performing models.

Custom training

Custom Training

Tensorflow Files that must be modified for custom dataset training

- train_utils.py
- data_generator.py
- Eval.py



core	Fixed improper HNASNet architecture (#6419)	5 months ago
datasets	Update download_and_convert_ade20k.sh (#7160)	last month
deprecated	Open source deeplab changes. Check change log in README.md for detail...	6 months ago
evaluation	Adding panoptic evaluation tools and update internal changes. (#6320)	5 months ago
g3doc	Adding quantization support for deeplab (#6681)	3 months ago
testing	Open source deeplab changes. Check change log in README.md for detail...	6 months ago
utils	Open source deeplab changes. Check change log in README.md for detail...	6 months ago
README.md	Adding panoptic evaluation tools and update internal changes. (#6320)	5 months ago
__init__.py	Added deeplab model to research folder	2 years ago
common.py	Open source deeplab changes. Check change log in README.md for detail...	6 months ago
common_test.py	PiperOrigin-RevId: 205684720	last year
deeplab_demo.ipynb	Update colab demo (#6652)	4 months ago
eval.py	Fix DeepLab Python 3 compatibility issue #6797. (#6799)	3 months ago
export_model.py	Adding quantization support for deeplab (#6681)	3 months ago
input_preprocess.py	Open source deeplab changes. Check change log in README.md for detail...	6 months ago
local_test.sh	Adding quantization support for deeplab (#6681)	3 months ago
local_test_mobilenetv2.sh	Adding quantization support for deeplab (#6681)	3 months ago
model.py	Open source deeplab changes. Check change log in README.md for detail...	6 months ago
model_test.py	Open source deeplab changes. Check change log in README.md for detail...	6 months ago
train.py	Fix DeepLab Python 3 compatibility issue #6797. (#6799)	3 months ago
vis.py	Fix DeepLab Python 3 compatibility issue #6797. (#6799)	3 months ago

Steps

1. Modify tensorflow files with your dataset information
2. Convert masks to color-indexed images
3. Create
 - a. 'train.txt' with training image filenames
 - b. 'val.txt' with validation image filenames
 - c. 'trainval.txt' with both
 - d. NOTE: training images and masks should have the same filename
4. Create a tfrecord of your dataset using 'build_voc2012_data.py'
 - a. Tfrecord = a Tensorflow binary storage format
5. If using a pre-trained model (transfer learning)
 - a. https://github.com/tensorflow/models/blob/master/research/deeplab/g3doc/model_zoo.md
6. Run 'train.py'
7. Run 'eval.py' - custom file has miou per class and overall miou
8. Run 'vis.py' - to get segmentation masks
9. Run 'export_model.py' to save your trained 'frozen_inference_graph.pb'
10. Use 'deeplab_demo.ipynb' with your checkpoint for future inference

Custom Training - Datasets


- Labeled Faces in the Wild - Parts Dataset
- CelebA segmentation masks & CelebA

Preprocessing needs to be done to convert segmentation maps to color-indexed images.


0 = background
1 = hair
2 = skin



LFW image



LFW - parts mask



color-indexed mask



Custom Training Results

Weights for classes:

1, background
10, hair
5, face


10000 iterations: miou

class_0: 0.91722
class_1: 0.66235
class_2: 0.83431
mean_iou: 0.8046266666666666



20000 iterations: miou

class_0: 0.92631
class_1: 0.67995
class_2: 0.85498
mean_iou: 0.8204133333333333



Bibtex format References

```
[1] @inproceedings{deeplabv3plus2018,
  title={Encoder-Decoder with Atrous Separable Convolution for Semantic Image Segmentation},
  author={Liang-Chieh Chen and Yukun Zhu and George Papandreou and Florian Schroff and Hartwig Adam},
  booktitle={ECCV},
  year={2018}
}

[2] @article{DBLP:journals/corr/ChenPSA17,
  author = {Liang{-}Chieh Chen and George Papandreou and Florian Schroff and Hartwig Adam},
  title = {Rethinking Atrous Convolution for Semantic Image Segmentation},
  journal = {CoRR},
  volume = {abs/1706.05587},
  year = {2017},
  url = {http://arxiv.org/abs/1706.05587},
  archivePrefix = {arXiv},
  eprint = {1706.05587},
  timestamp = {Mon, 13 Aug 2018 16:48:07 +0200},
  biburl = {https://dblp.org/rec/bib/journals/corr/ChenPSA17},
  bibsource = {dblp computer science bibliography, https://dblp.org}
}

[3] @article{DBLP:journals/corr/ChenPK0Y16,
  author = {Liang{-}Chieh Chen and George Papandreou and Iasonas Kokkinos and Kevin Murphy and Alan L. Yuille},
  title = {DeepLab: Semantic Image Segmentation with Deep Convolutional Nets,
    Atrous Convolution, and Fully Connected CRFs},
  journal = {CoRR},
  volume = {abs/1606.00915},
  year = {2016},
  url = {http://arxiv.org/abs/1606.00915},
  archivePrefix = {arXiv},
  eprint = {1606.00915},
  timestamp = {Mon, 13 Aug 2018 16:46:13 +0200},
  biburl = {https://dblp.org/rec/bib/journals/corr/ChenPK0Y16},
  bibsource = {dblp computer science bibliography, https://dblp.org}
}
```